

Safety


HS-Guard v10

Qwen3Guard 0.6B

面向生成过程的安全审核。固定工作点下，同时提升红线回答识别与正常回答放行。


Response Recall

回答红线召回率 (%) ↑ · 114 examples




False Positives

安全回答误报率 (%) ↓ · 203 examples




ToxicChat

Prompt F1 (%) ↑ · General safety



Aegis 2.0

Response F1 (%) ↑ · General safety



+6.14 pp recall

33.3% fewer false positives

Same-example evaluation. Benchmark definitions and operating points: Technical Report.

Streaming

HS-Guard v10

Qwen3Guard 0.6B

单卡实时审核。历史上下文从 512 扩展到 8,000 tokens，增量服务延迟保持约 2.38 ms。

512 tokens


Single-token latency, P50 (ms) ↓



4.9x faster

Concurrent Moderation

NVIDIA L20 · 30-60 tokens/s per stream



P95 <= 20 ms at all three contexts

2,048 tokens

Single-token latency, P50 (ms) ↓



4.9x faster

8,000 tokens

Single-token latency, P50 (ms) ↓



5.3x faster



Minimum across six measurements



Worst of two seeds at 352 streams

Warm incremental serving. Qwen baseline: SGLang 0.5.3rc0 with radix cache enabled.

Model-engine timings; tokenization, prefill and transport are outside the measured interval.

Open Release

模型、流式推理引擎和技术报告。

 HS-Guard v10

 Qwen3Guard 0.6B

HS-Guard v10

754.5M parameters · 24 layers · Hybrid recurrent state + W512 attention

分别提供提问与回答审核，以及红线和通用安全两种分类分支。

只计算新增 token 的风险分数，可接入模型生成过程。

GitHub

流式引擎、加载器与示例

<https://github.com/KrisLiu16/HS-Guard>

Hugging Face

完整权重与 tokenizer

<https://huggingface.co/KrisLiu16/HS-Guard-v10>

Technical Report

架构、训练与评测方法

https://github.com/KrisLiu16/HS-Guard/blob/main/docs/technical_report.pdf

Code: Apache-2.0 | Weights: CC BY-NC 4.0 (noncommercial research)

Zhihao Liu · Naive N0.5